Generalization

Bert Huang
Data Intensive Studies Center
Tufts University

Outline

- Introduction
- Bias and Variance (interactive website)
- Cross-Validation
- Relationship to my research

Complexity in Machine Learning

- Relates to issues in philosophy of science, induction
  - Occam's Razor, etc.
- Baseball (sports) statistics

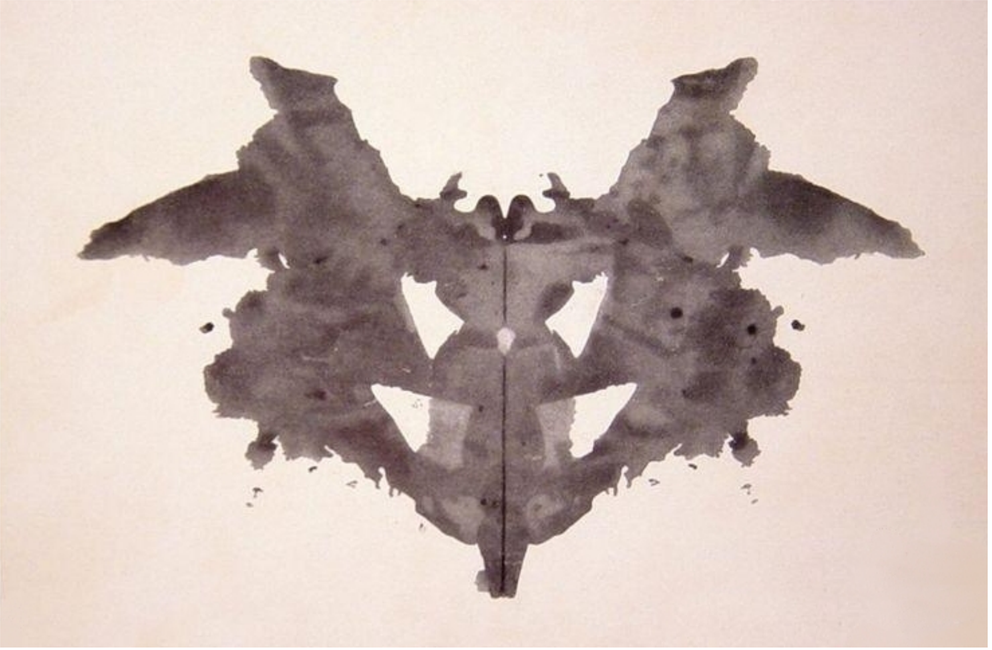

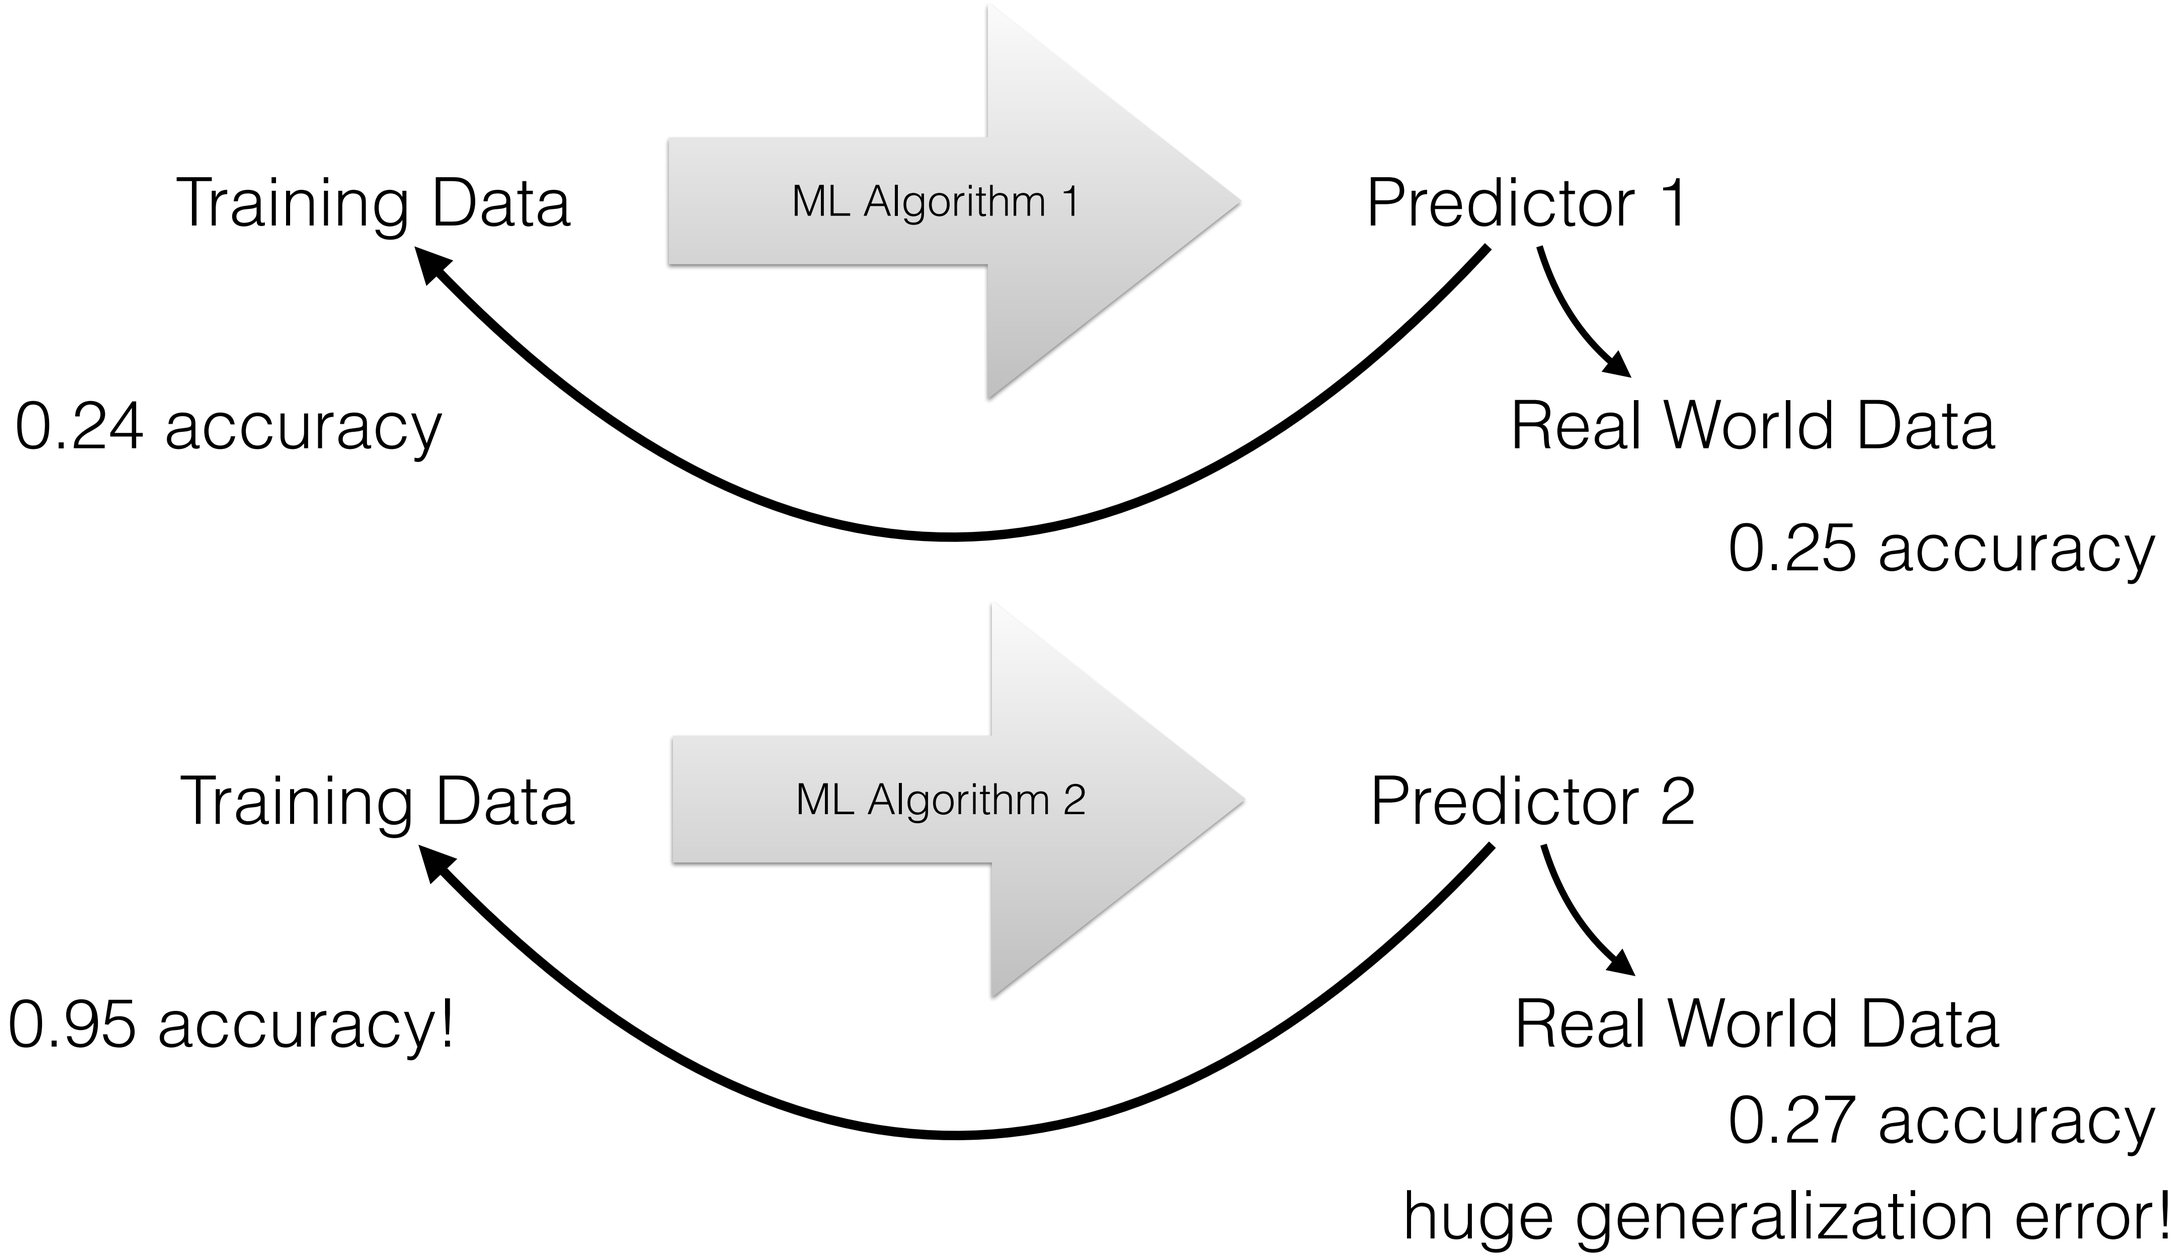
Training Data
ML Algorithm 1
Predictor 1
0.24 accuracy
Real World Data
0.25 accuracy
Training Data
ML Algorithm 2
Predictor 2
0.95 accuracy!
Real World Data
0.27 accuracy
huge generalization error!

Underfitting

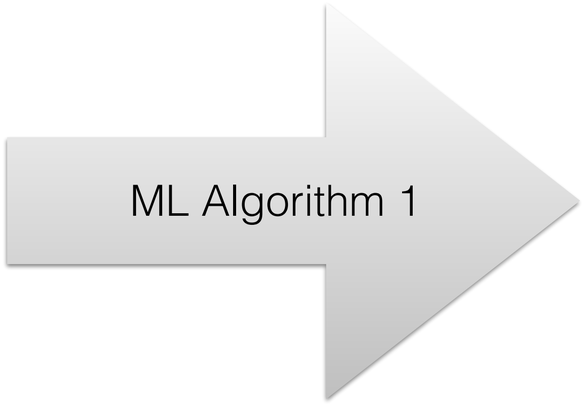

- Low dimensional
- Heavily regularized
- Bad modeling assumptions

Overfitting

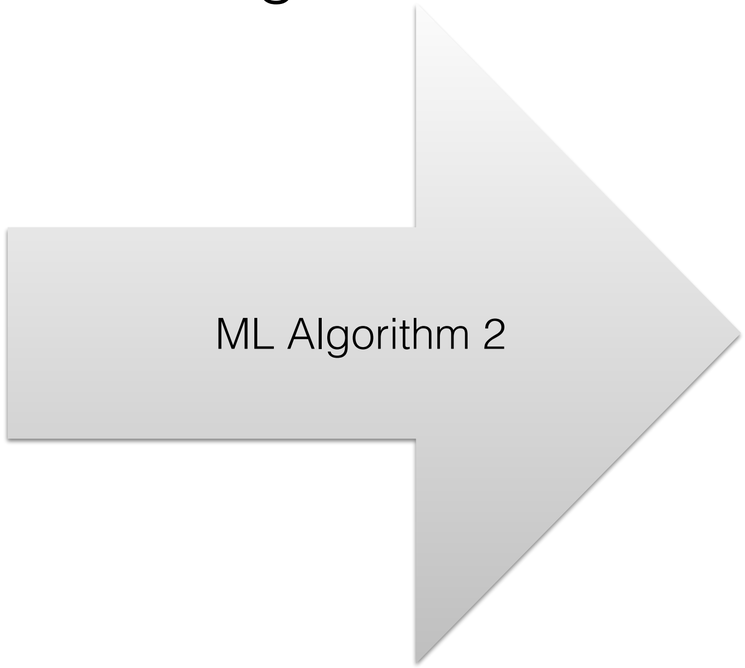

- High dimensional or non-parametric
- Weakly regularized
- Not enough modeling assumptions
- Not enough data

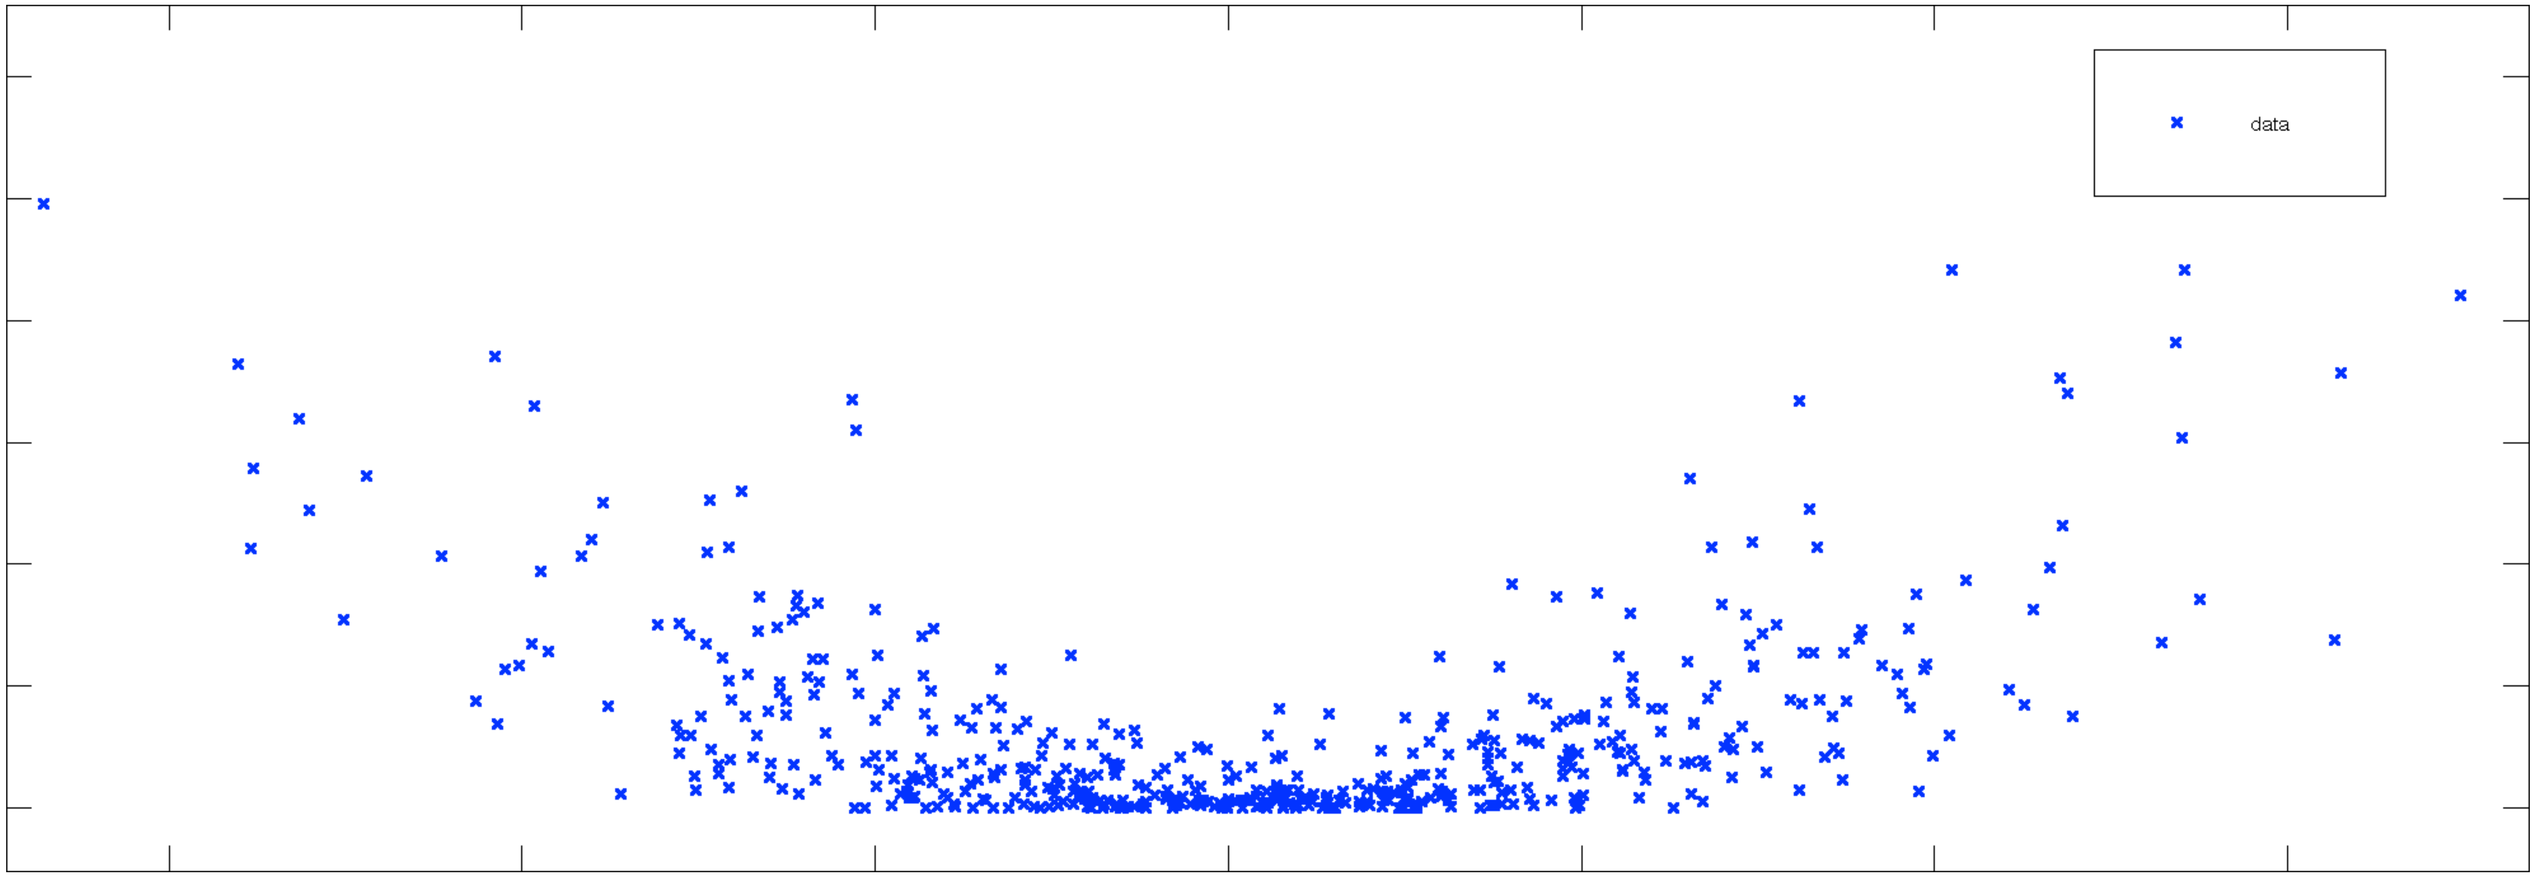
data

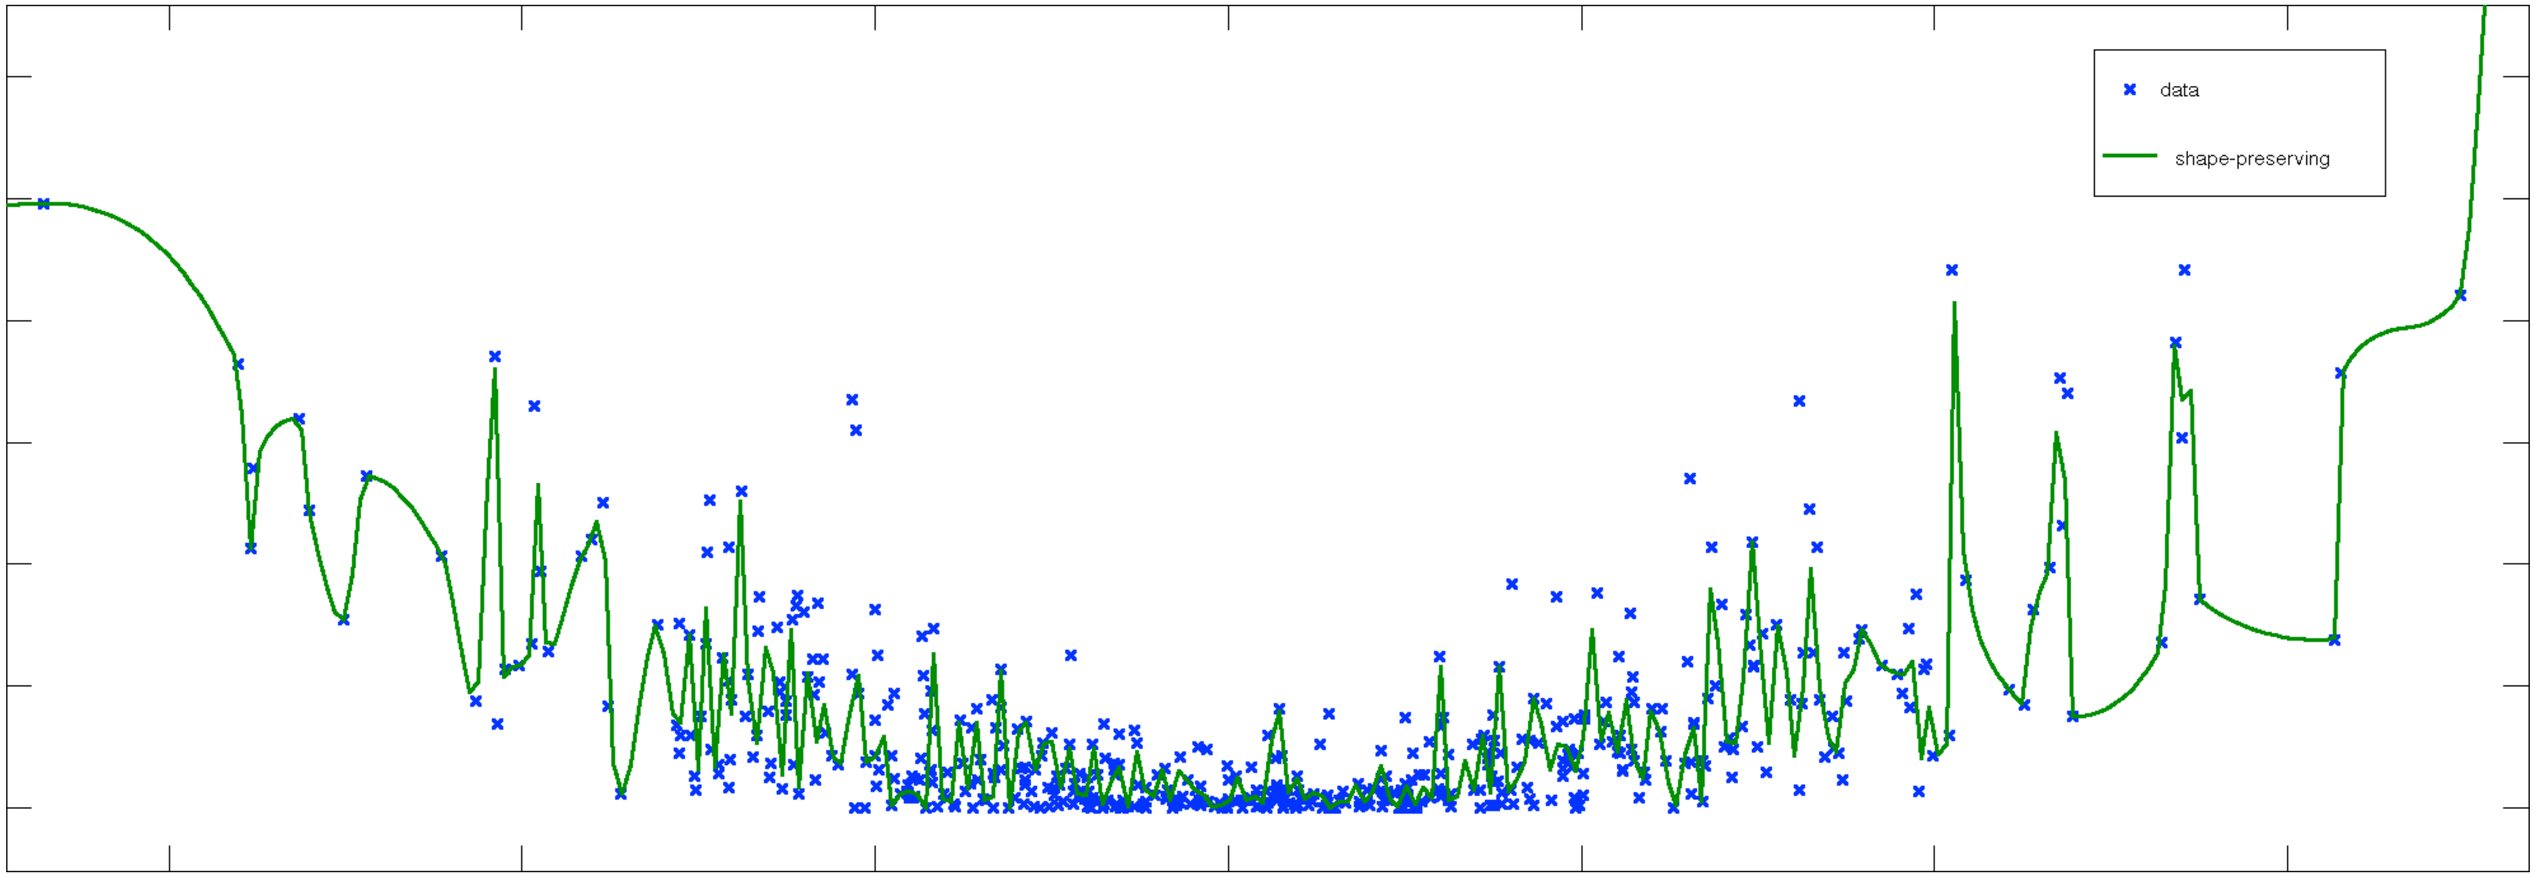
data
shape-preserving

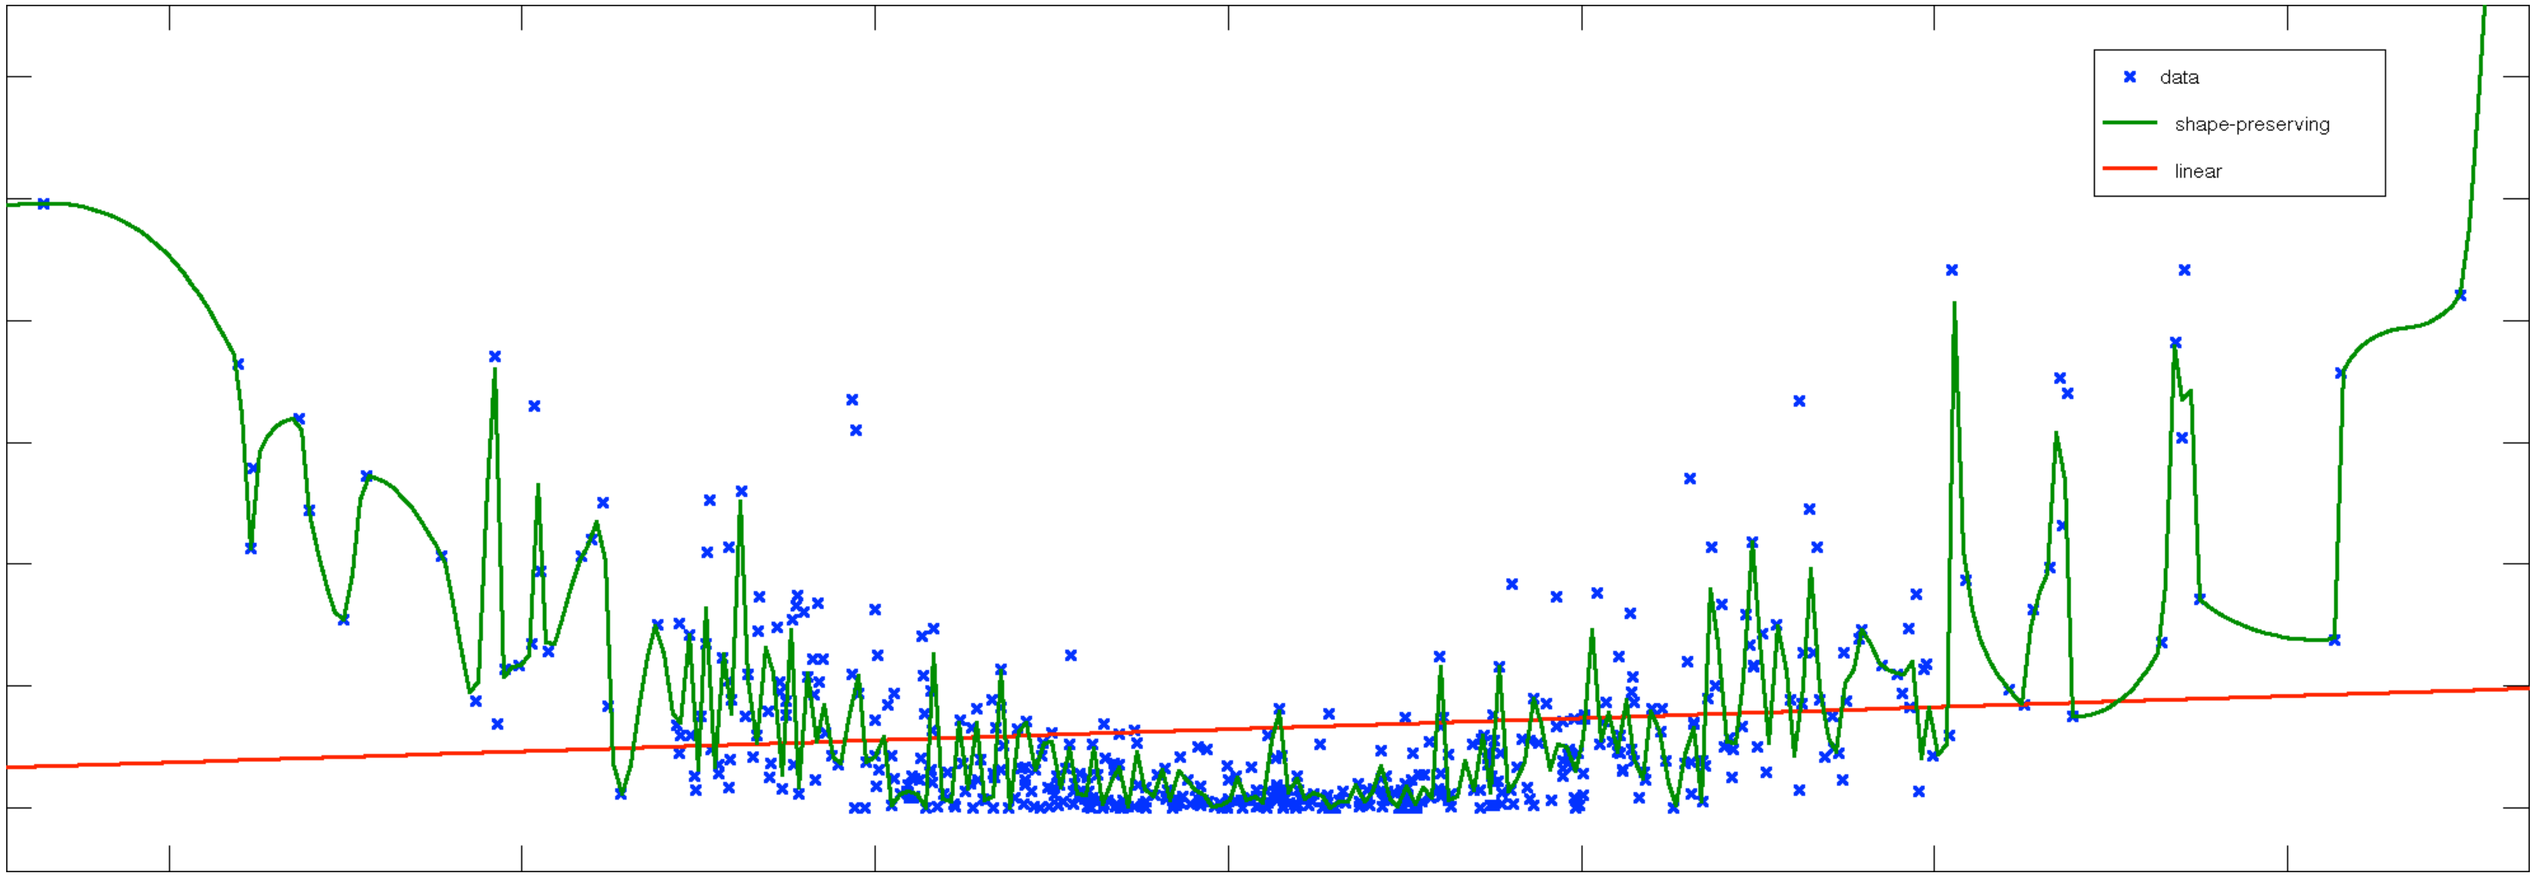
data
shape-preserving
linear

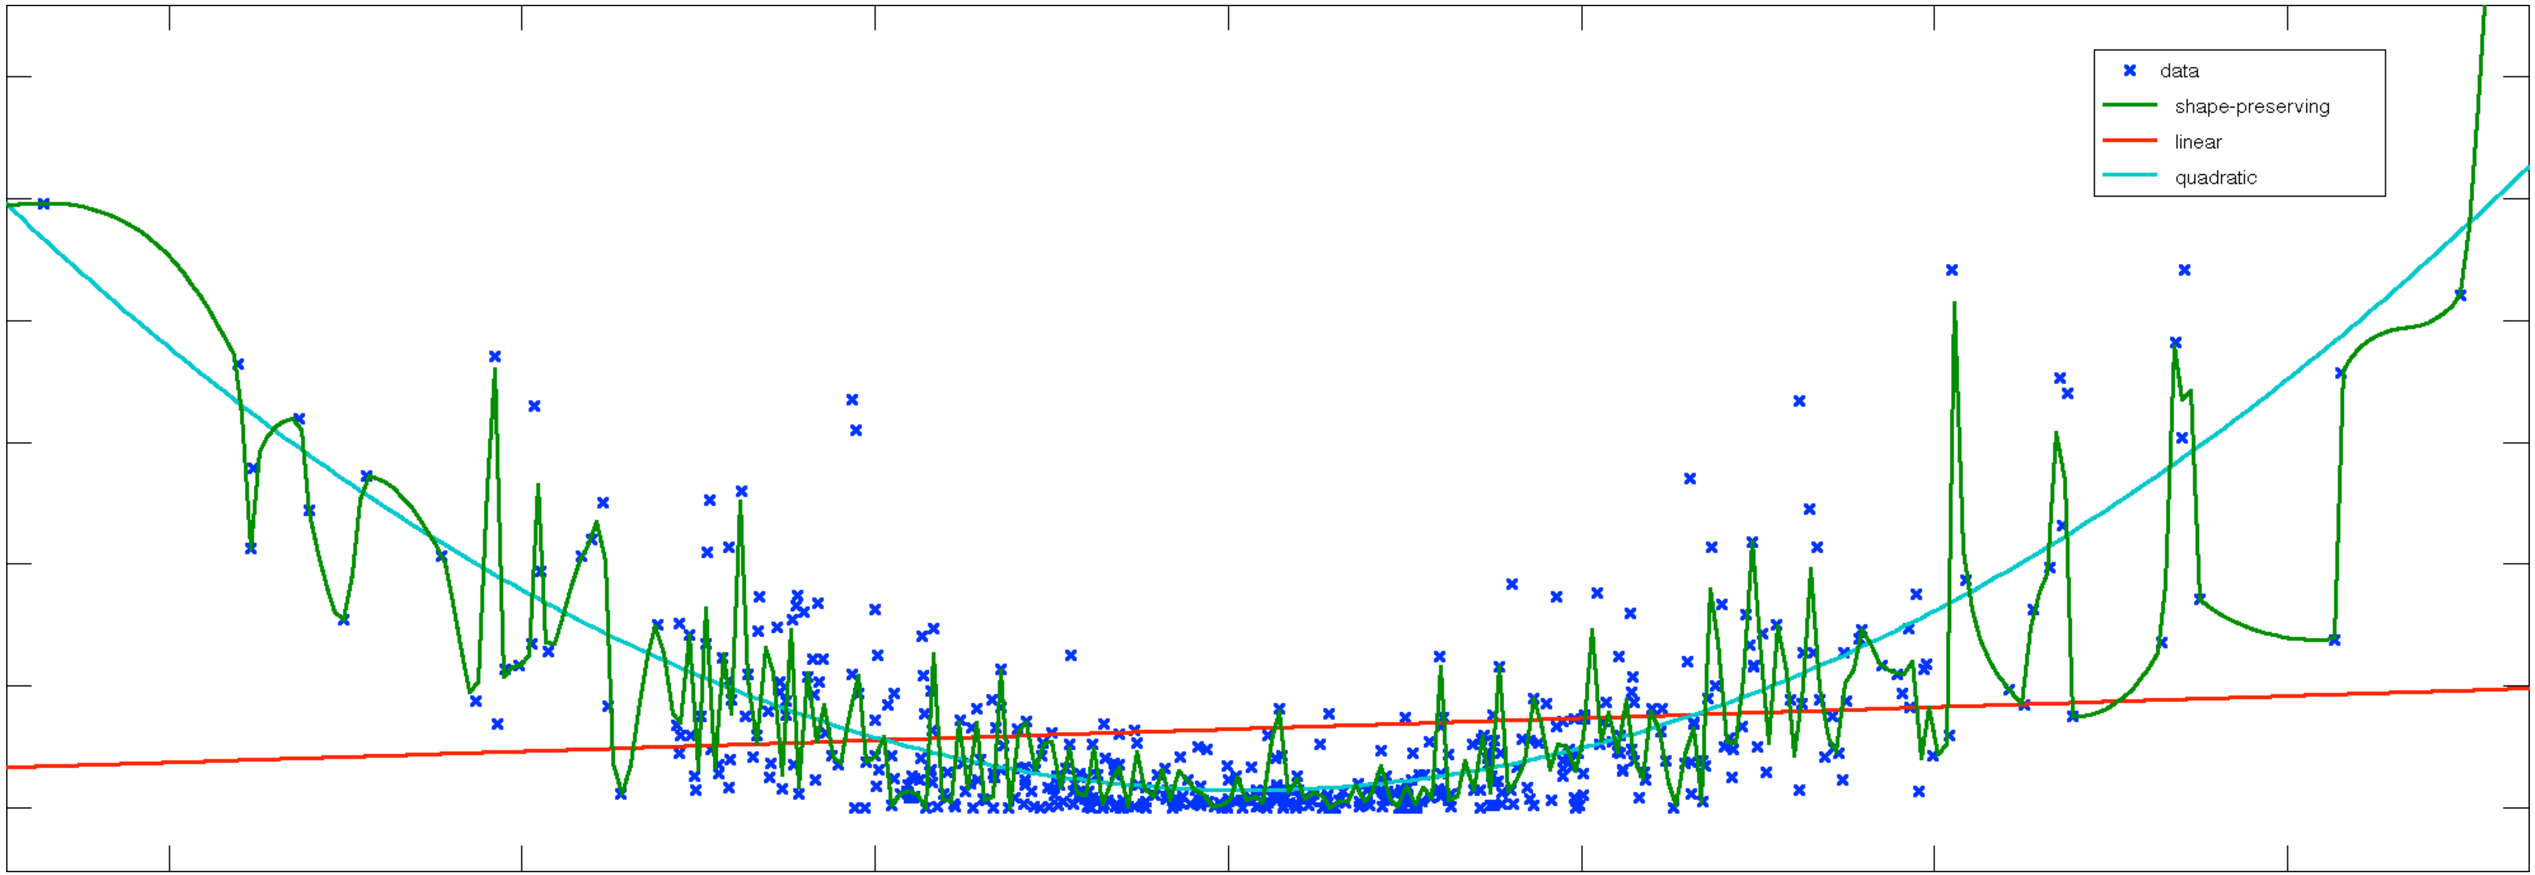
data
shape-preserving
linear
quadratic

Overfitting and Underfitting

- Training models too complex can cause overfitting
- Training models too simple (or wrong) can cause underfitting

Outline

- Overfitting and underfitting
- Bias and variance
- Validation for model selection

http://www.r2d3.us/visual-intro-to-machine-learning-part-2/

Nearest-Neighbor Classifiers

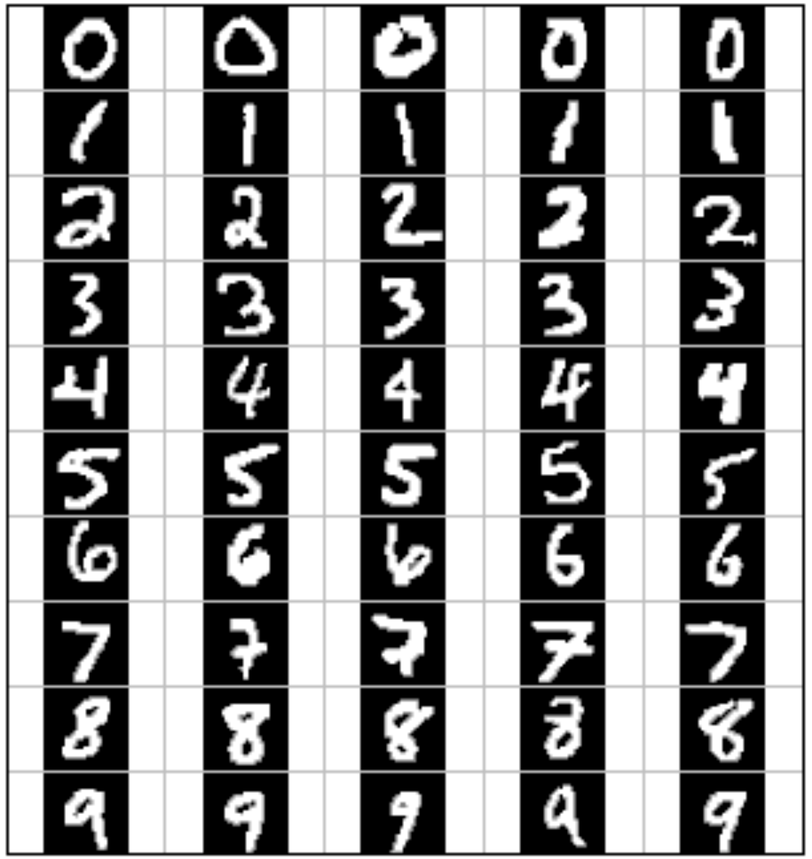

classifier = {

 : 0,

 : 0,

 : 0,

 : 0,

 : 0,

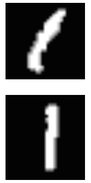 : 1,

 : 1,

…

}

100% training accuracy!

53% testing accuracy…

Held-out Validation

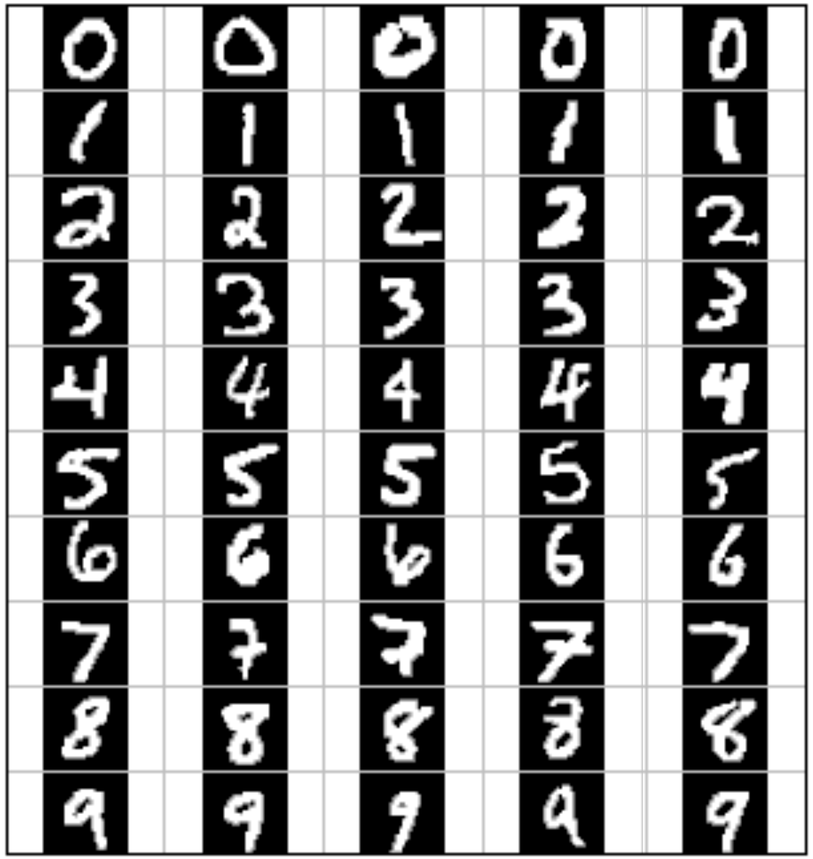

Held-out Validation

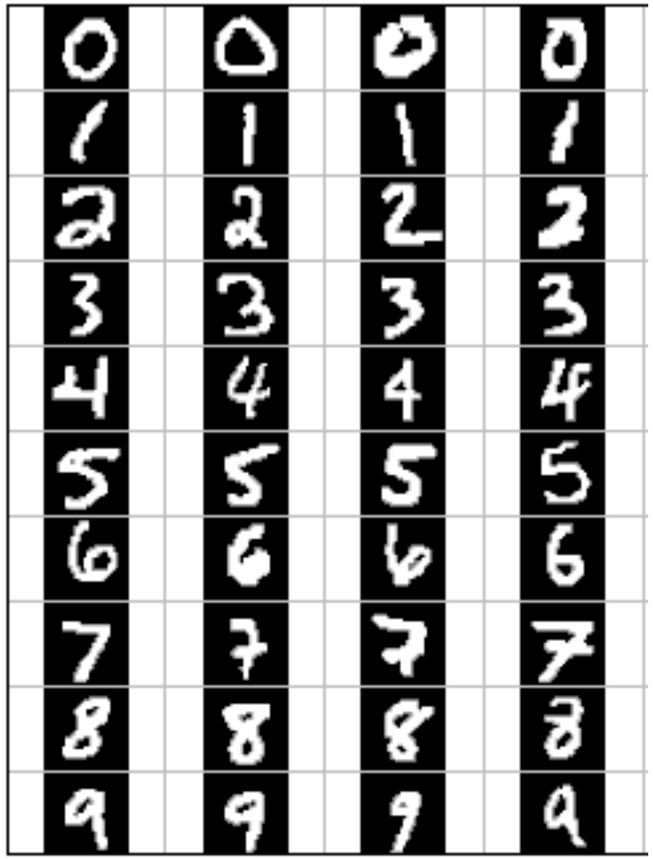

training data

	Accuracy on training data	Accuracy on validation data
Simple	0.91	0.83
Medium	0.95	0.88
Complex	0.99	0.79
Super Complex	1.0	0.54

validation data

Cross Validation

Fold 1

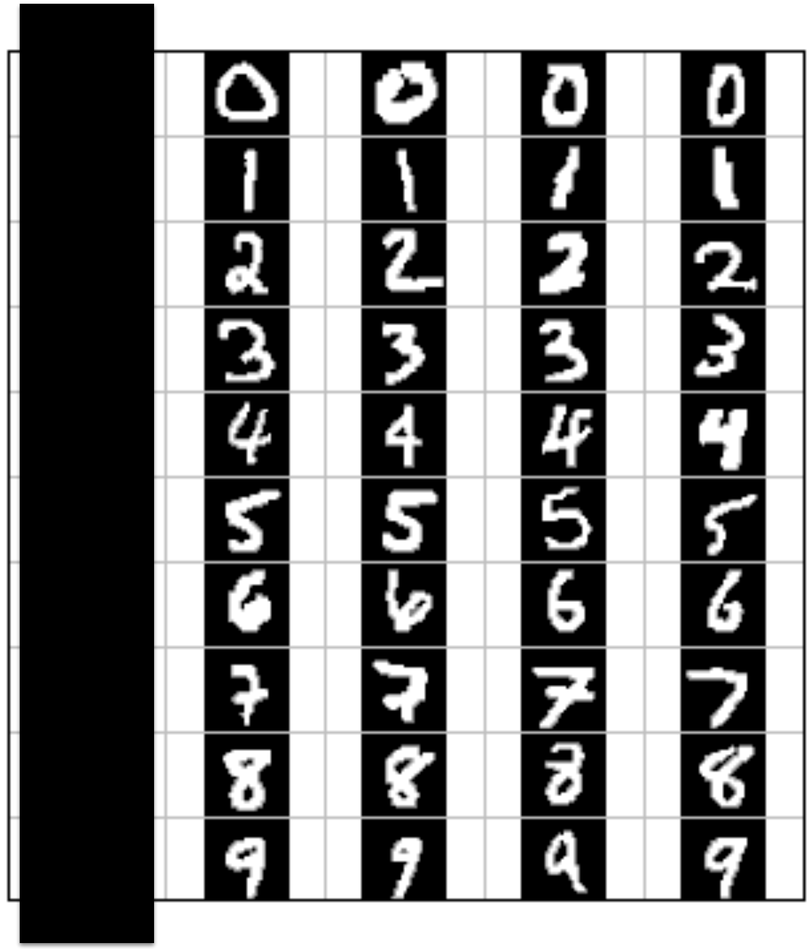

training data

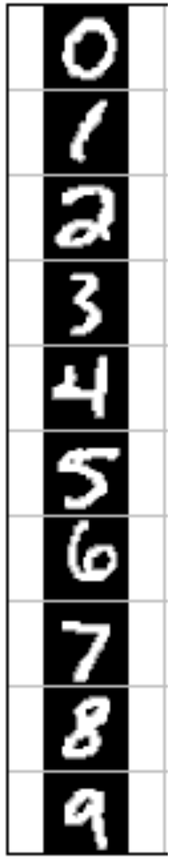

validation data

Cross Validation

Fold 2

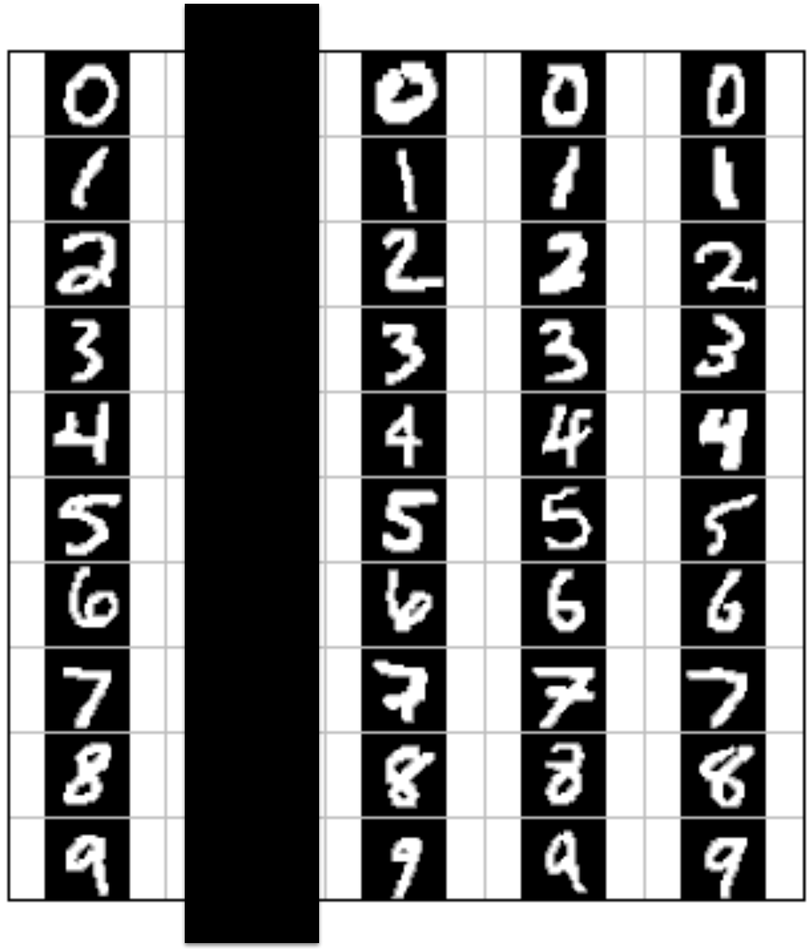

training data

validation data

Cross Validation

Fold 3

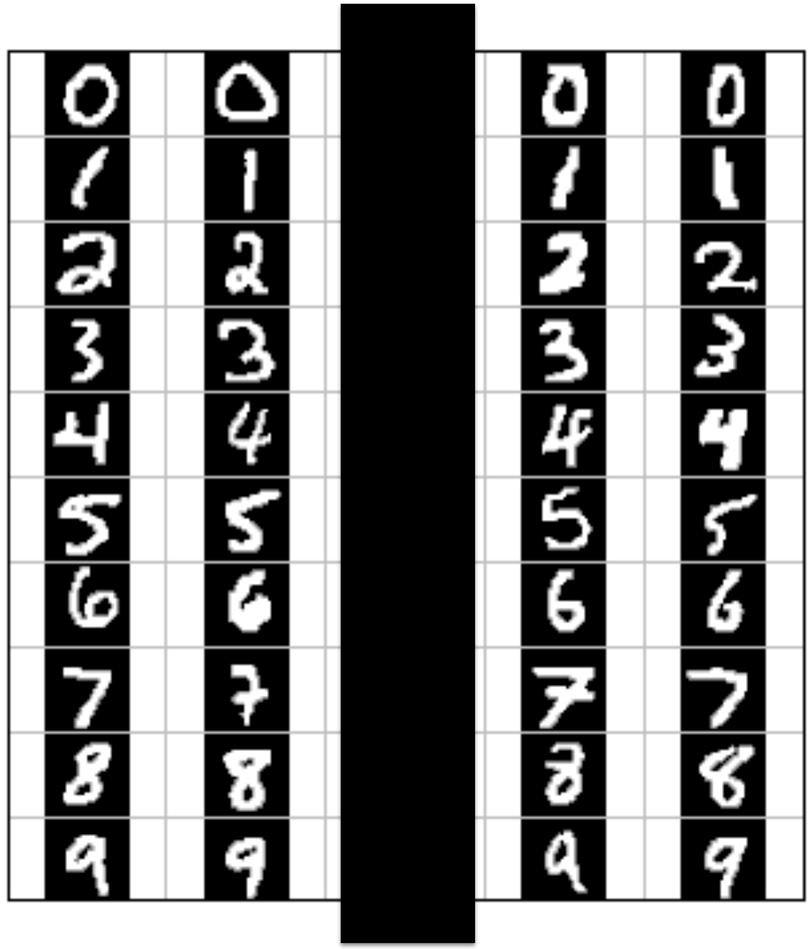

training data

validation data

Cross Validation

Fold 4

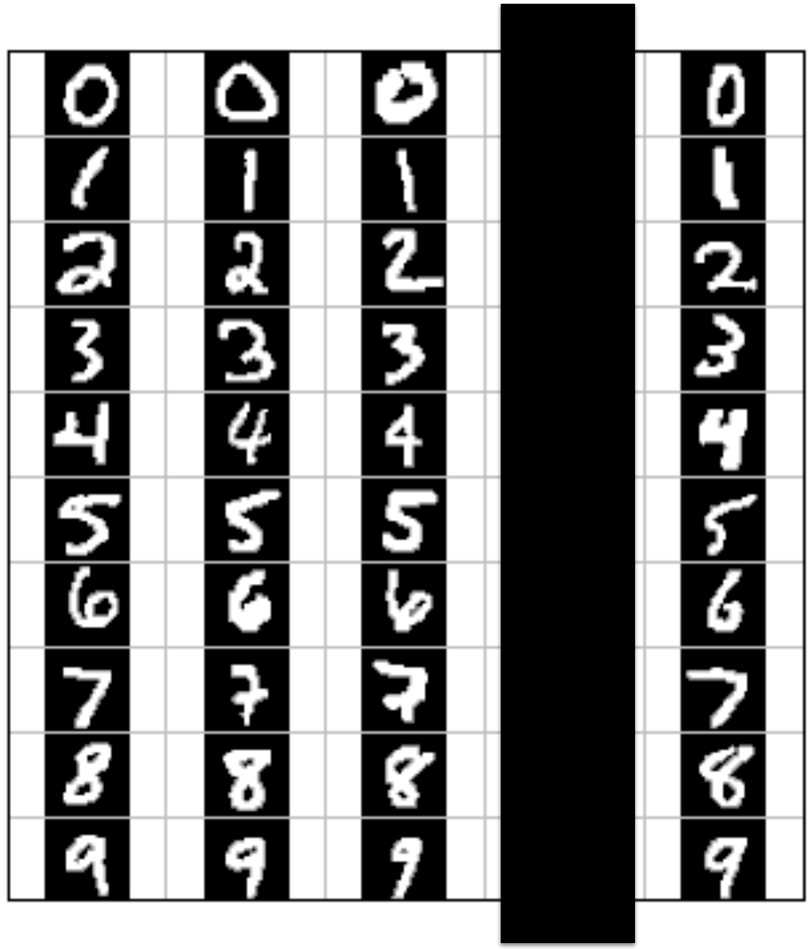

training data

validation data

Cross Validation

Fold 5

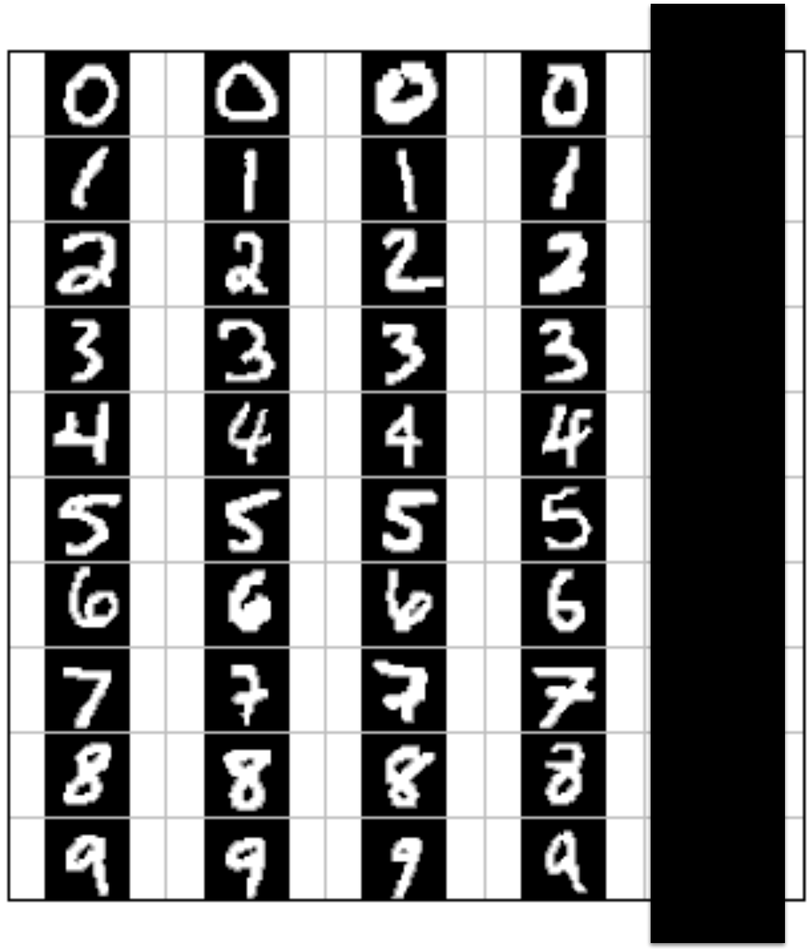

training data

validation data

Leave-one-out Cross Validation

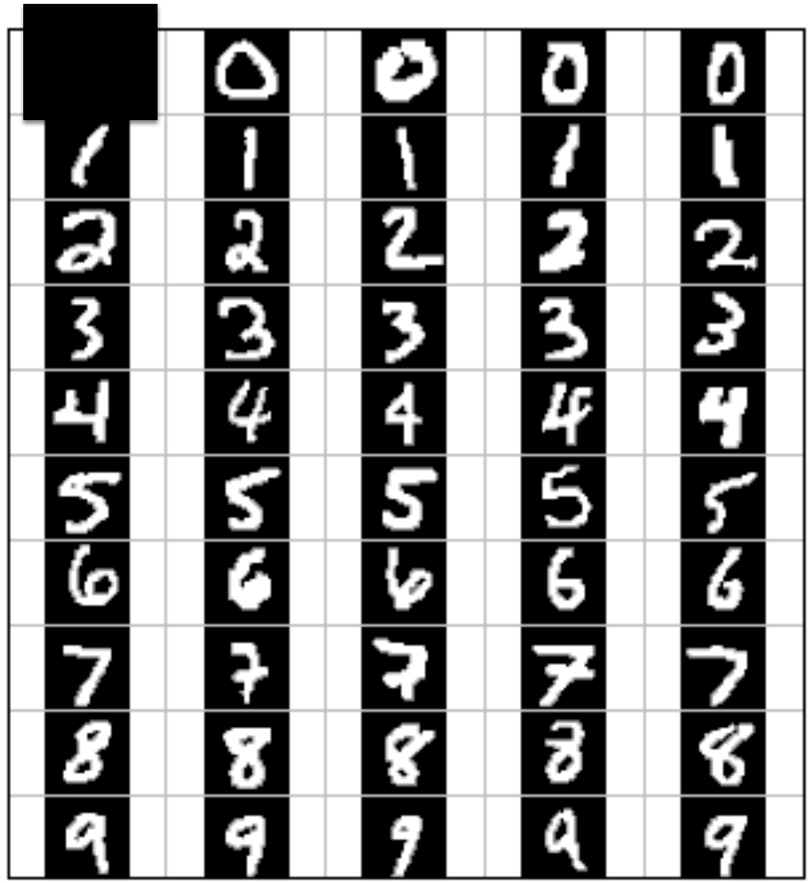

training data

validation data

Leave-one-out Cross Validation

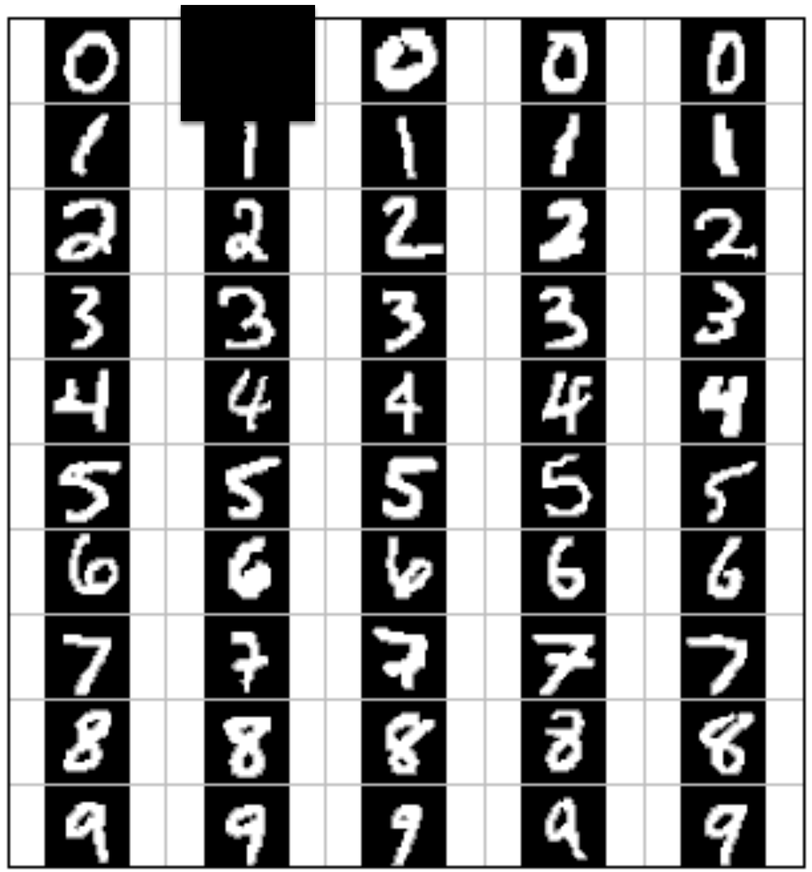

training data

validation data

Leave-one-out Cross Validation

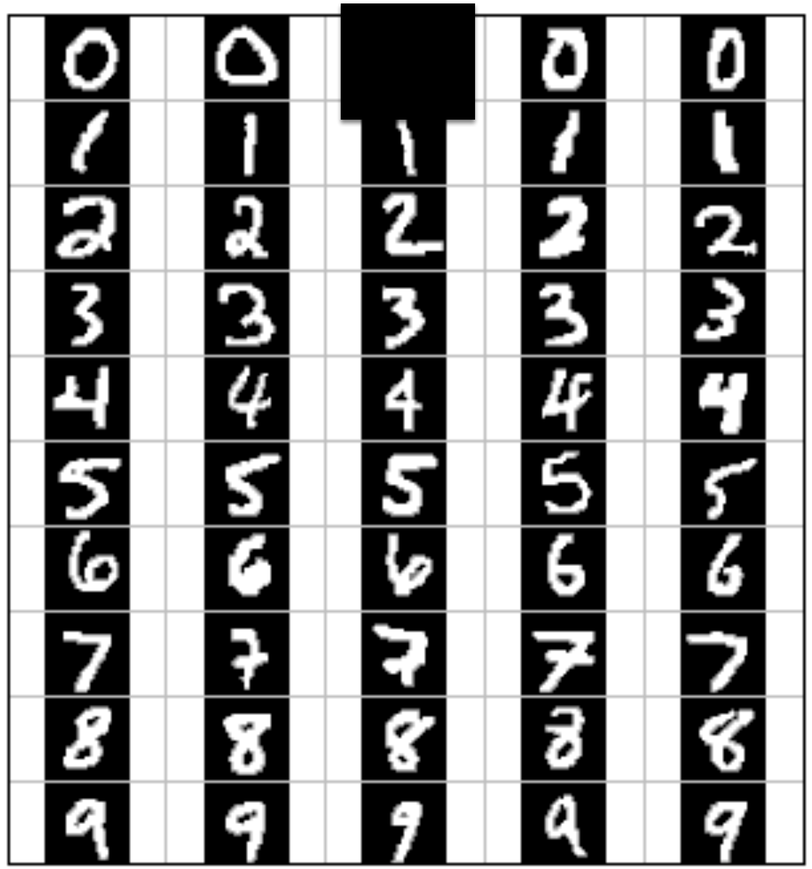

training data

validation data

Leave-one-out Cross Validation

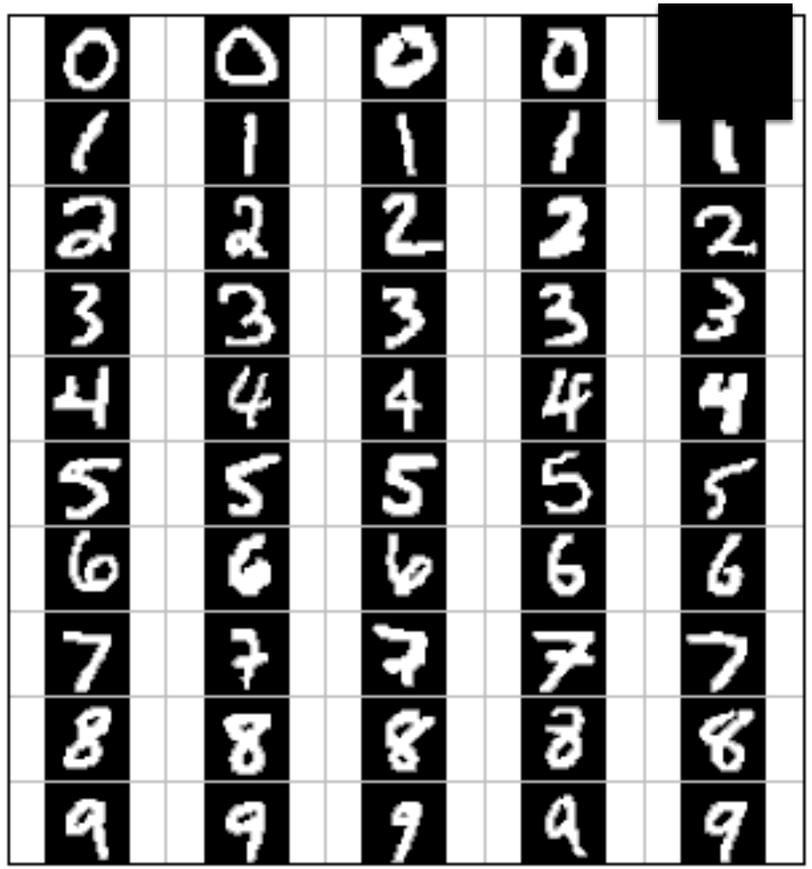

training data

validation data

Leave-one-out Cross Validation

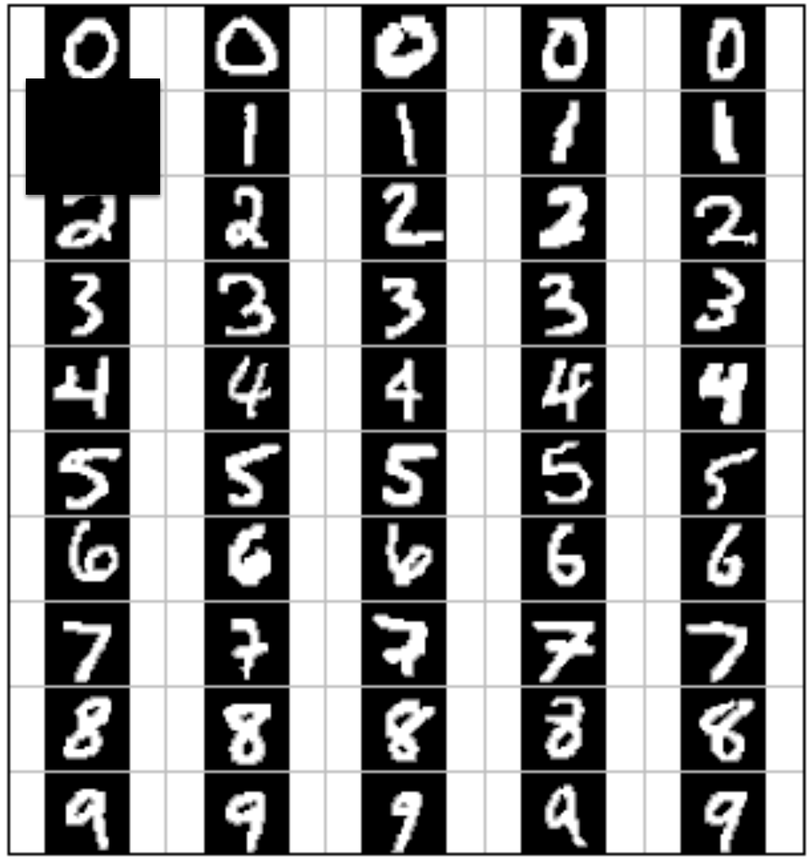

training data

validation data

Leave-one-out Cross Validation

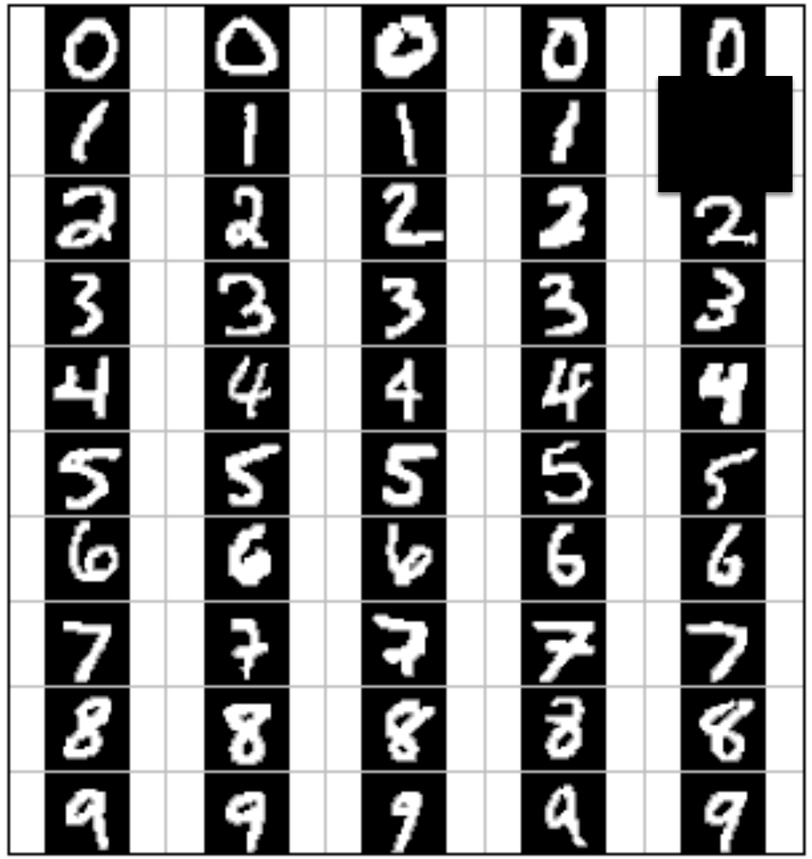

training data

validation data

Leave-one-out Cross Validation

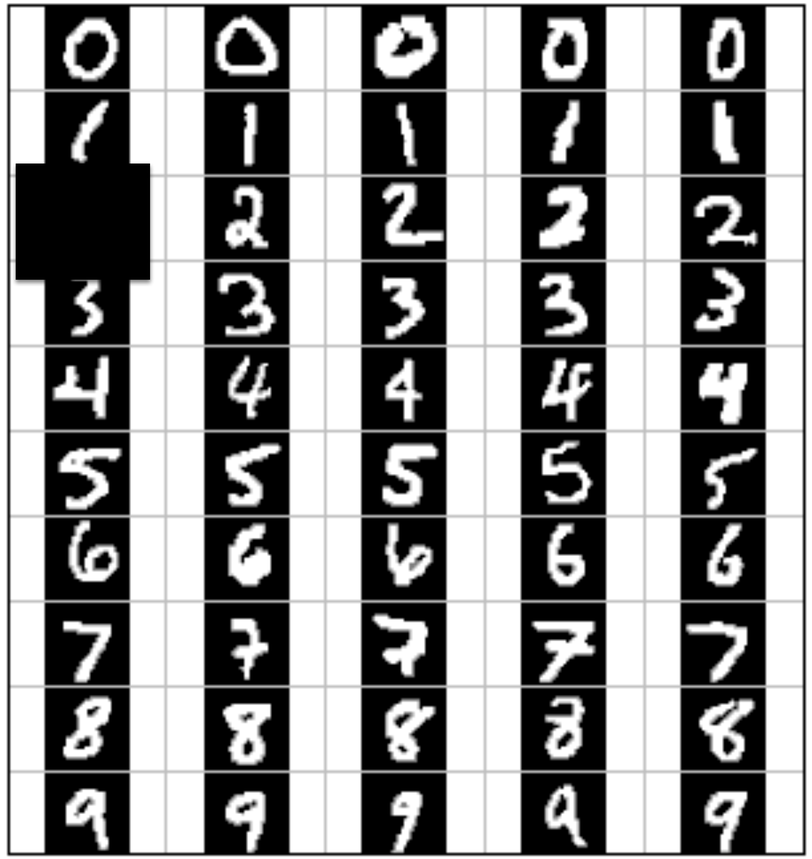

training data

validation data

How Many Folds?

- What are the pros and cons of leave-one-out cross-validation?
- We usually train on N-1 folds and test on 1 fold. What are pros and cons of doing the inverse: train on 1 fold and test on N-1 folds?

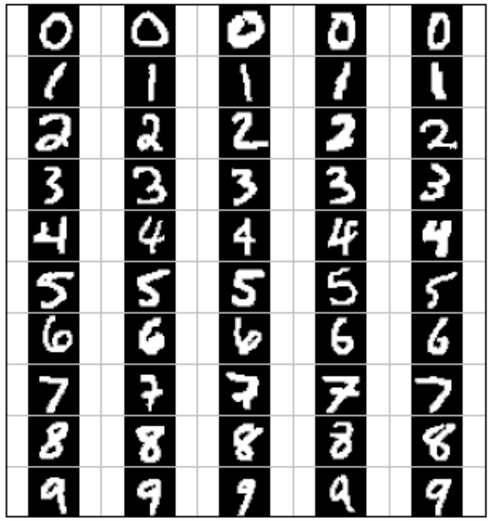

Training

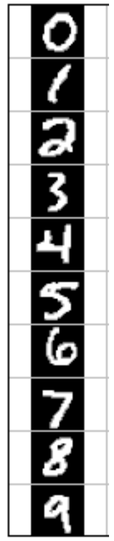

Testing

How Many Folds?

- What are the pros and cons of leave-one-out cross-validation?
- We usually train on N-1 folds and test on 1 fold. What are pros and cons of doing the inverse: train on 1 fold and test on N-1 folds?

Training

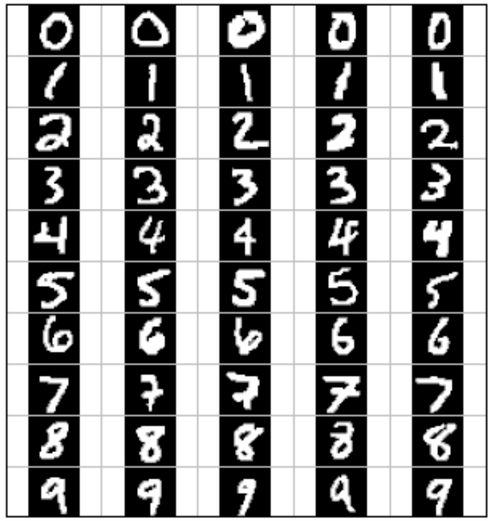

Testing

Testing versus Validation

- Best practice for experiments:
  - Hold out test set completely hidden from training
  - Use validation on training data for model (or parameter) selection
  - Evaluate on held-out test data

Model Selection via Validation

- Measure performance on **held-out** training data
  - Simulate testing environment
- Rotate **folds** of held-out subsets
- Can even hold out one at a time: **leave-one-out** validation
- Use (cross) validation performance to tune extra parameters

Take-Away Points

- Overfitting and underfitting, bias and variance
  - bias -> modeling error, variance -> sampling error
  - Always have a mix of two
- Validation for model selection
- Reducing bias may need more complex models, but comes with challenges that researchers are working on solving